

COMP9414: Artificial Intelligence

Lecture 2a: Problem Solving

Wayne Wobcke

e-mail:w.wobcke@unsw.edu.au

UNSW

© W. Wobcke et al. 2019

COMP9414

Problem Solving

1

This Lecture

- Search as a “weak method” of problem solving with wide applicability
- Uninformed search methods (use no problem-specific information)
- Informed search methods (use **heuristics** to improve efficiency)

启发法

UNSW

© W. Wobcke et al. 2019

Motivating Example

- You are in Romania on holiday, in Arad, and need to get to Bucharest
- What more information do you need to solve this problem?
- Once you have this information, how do you solve the problem?
- How do you know your solution is any good? What extra information would you need in order to evaluate the quality of your solution?

UNSW

© W. Wobcke et al. 2019

COMP9414

Problem Solving

3

State Space Search Problems

- **State space** — set of all states reachable from initial state(s) by any action sequence
- **Initial state(s)** — element(s) of the state space
- **Transitions**
 - ▶ **Operators** — set of possible actions at agent’s disposal; describe state reached after performing action in current state, **or**
 - ▶ **Successor function** — $s(x)$ = set of states reachable from state x by performing a single action
- **Goal state(s)** — element(s) of the state space
- **Path cost** — cost of a sequence of transitions used to evaluate solutions (apply to optimization problems)

UNSW

© W. Wobcke et al. 2019

Example Problem — 8-Puzzle

1	2	3
4	5	6
7	8	

States: location of eight tiles plus location of blank

Operators: move blank left, right, up, down

Goal state: state with tiles arranged in sequence

Path cost: each step is of cost 1

Real World Problems

- Route finding — robot navigation, airline travel planning, computer/phone networks
- Travelling salesman problem — planning movement of automatic circuit board drills
- VLSI layout — design ^硅silicon chips
- Assembly sequencing — scheduling assembly of complex objects, manufacturing process control
- Mixed/constrained problems — courier delivery, product distribution, fault service and repair

These are **optimization** problems but mathematical (operations research) techniques are not always effective.

Example Problem — N-Queens

						♛
	♛					
			♛			
♛						
					♛	
				♛		
		♛				
				♛		

States: 0 to N queens arranged on chess board

Operators: place queen on empty square

Goal state: N queens on chess board, none attacked

Path cost: zero

Problem Representation — Tic-Tac-Toe

X	O	X
X	O	O
X		O

States: arrangement of Os and Xs on 3x3 grid

Operators: place X (O) in empty square

Goal state: three Xs (Os) in a row

Path cost: zero

Tic-Tac-Toe — First Attempt

1	2	3
4	5	6
7	8	9

Board: 0=blank; 1=X; 2=O

Idea: Use move table with $3^9 = 19683$ elements 三元的

Algorithm: Consider board to be a ternary number; convert to decimal; access move table; update board

- Fast; lots of memory; laborious; not extensible

Tic-Tac-Toe — Third Attempt

8	3	4
1	5	9
6	7	2

Board is a magic square!

Algorithm: As in attempt 2 but to check for win — keep track of player's "squares". If difference of 15 and sum of two squares is ≤ 0 or > 9 two squares are not collinear. Otherwise, if square equal to difference is blank, move there.

- What does this tell you about the way humans solve problems vs computers?

Tic-Tac-Toe — Second Attempt

1	2	3
4	5	6
7	8	9

Board: 2=blank; 3=X; 5=O

Algorithm: Separate strategy for each move.

Goal test (if row gives win on next move): calculate product of values

X: test product = 18 ($3 \times 3 \times 2$); O: test product = 50 ($5 \times 5 \times 2$)

- Not as fast as 1; much less memory; easier to understand and comprehend; strategy determined in advance; not extensible

Tic-Tac-Toe — Fourth Attempt

	?	

Board: list of board positions arising from next move; estimate of likelihood of position leading to a win


Algorithm: look at position arising from each move; choose "best" one

- Slower; can handle large variety of problems

Back to Motivating Example

- Notice assumptions built in to problem formulation (level of abstraction)
- Note that while people can “look” at the map to see a solution, the computer must construct the map by exploration
 - Where can I go from Arad?
 - Sibiu, Timisoara, Zerind
 - Where can I go from Sibiu?
- The order of questioning defines the search strategy
- Problem formulation assumptions critically affect the quality of the solution to the original problem

State Space — 8-Puzzle



Explicit State Spaces

- View state space search in terms of finding a path through a graph
- Graph** $G = (V, E)$ — V : vertices; E : edges
- Edges may have associated **cost**; **path cost** = sum edge costs in path
- Path** from vertex s to g — sequence of vertices $s = n_0, n_1, \dots, n_k = g$ such that there is an edge from n_i to n_{i+1}
- State space graph** — node represents state; edge represents change from one state to another due to action; costs may be associated with vertices and edges (hence paths)
- Forward (backward) branching factor** — max #out-(in-)going arcs from (to) node

Complications

- Single-state — agent starts in known world state and knows which unique state it will be in after a given action
- Multiple-state — limited access to world state means agent is unsure of world state but may be able to narrow it down to a set of states
- 偶发事件** **Contingency** problem — if agent does not know full effects of actions (or there are other things going on) it may have to sense during execution (changing the search space dynamically)
- Exploration problem — no knowledge of effects of actions (or state), so agent must experiment

Search methods are capable of tackling single-state problems and multiple-state problems at the cost of additional complexity

Uninformed (Blind) Search Algorithms

- Breadth-First Search
- Uniform Cost Search
- Depth-First Search
- Depth-Limited Search
- Iterative Deepening Search
- Bidirectional Search


General Search Procedure

```

function GeneralSearch(problem, strategy) returns a solution or failure
  initialize search graph using the initial state of problem
  loop
    if there are no candidates for expansion then return failure
    choose a frontier node for expansion according to strategy
    if the node contains a goal state then return solution
    else expand the node and add the resulting nodes to the search graph
  end
  
```

Note: Only test whether at goal state when expanding node, not when adding nodes to the search graph (except for breadth-first search!)


General Search Space (not State Space)



Search strategy — way in which frontier expands

State Space vs Search Space

- A **state** is part of the formulation of the search problem
- A **node** is a data structure used in a search graph/tree, and includes:
 - ▶ **parent**, **operator**, **depth**, **path cost** $g(x)$
- **States** do not have parents, children, depth, or path cost!



Two different nodes can have the same state

Evaluating Search Algorithms

- **Completeness:** strategy guaranteed to find a solution when one exists?
- **Time complexity:** how long to find a solution?
- **Space complexity:** memory required during search?
- **Optimality:** when several solutions exist, does it find the “best”?

Note: States are constructed during search, not computed in advance, so efficiently computing **successor states** is critical!

继承者 (后续的) (状态)

Breadth-First Search

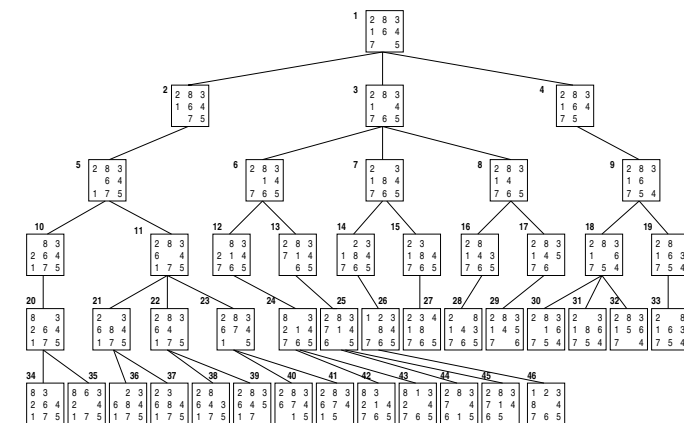
- **Idea:** Expand root node, then expand all children of root, then expand their children, ...
- All nodes at depth d are expanded before nodes at $d + 1$
- Can be implemented by using a **queue** to store frontier nodes
- Breadth-first search finds shallowest goal state
- **Stop when node with goal state is generated**
- **Include check that generated state has not already been explored**
 - ▶ Needs a new data structure for set of explored **states**

Analysis of Algorithms – Big-O

- $T(n)$ is $O(f(n))$ means that there is some n_0 and k such that

$$T(n) \leq kf(n) \text{ for every problem of size } n \geq n_0$$
- Independent of implementation, compiler, fixed overheads, ...
- $O()$ abstracts over constant factors
- Examples
 - ▶ $O(n)$ algorithm is better than an $O(n^2)$ algorithm (in the long run)
 - ▶ $T(100 \cdot n + 1000)$ is better than $T(n^2 + 1)$ for $n > 110$
 - ▶ Polynomial $O(n^k)$ **much** better than exponential $O(2^n)$
- $O()$ notation is a compromise between precision and ease of analysis

Breadth-First Search




Breadth-First Search — Analysis

- Complete
- Optimal — provided path cost is nondecreasing function of the depth of the node
- Maximum number of nodes generated: $b + b^2 + b^3 + \dots + b^d$ (where b = forward branching factor; d = path length to solution)
- Time and space requirements are the same $O(b^d)$

Uniform Cost Search

Uniform cost search is optimal only if it stops when goal node is **expanded** – not when goal node is generated



Uniform Cost Search

- Also known as **Lowest-Cost-First** search
- Shallowest goal state may not be the least-cost solution
- **Idea:** Expand lowest cost (measured by path cost $g(n)$) node
- Order nodes in the frontier in increasing order of path cost
- Breadth-first search \approx uniform cost search where $g(n) = \text{depth}(n)$ (except breadth-first search stops when goal state generated)
- Include check that generated state has not already been explored
- Include test to ensure frontier contains only one node for any state – for path with lowest cost

Uniform Cost Search — Analysis

- Complete
- Optimal — provided path cost does not decrease along path (i.e. $g(\text{successor}(n)) \geq g(n)$ for all n)
- Reasonable assumption when path cost is cost of applying operators along the path
- Performs like breadth-first search when $g(n) = \text{depth}(n)$
- If there are paths with negative cost, need exhaustive search

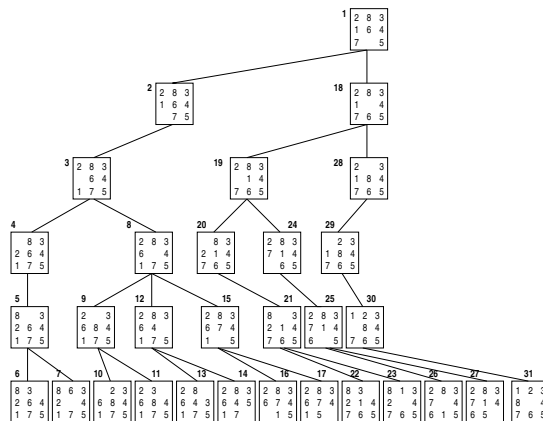
Depth-First Search

- **Idea:** Always expand node at deepest level of tree and when search hits a dead-end return back to expand nodes at a shallower level
- Can be implemented using a **stack** of explored + frontier nodes
- At any point depth-first search stores single path from root to leaf together with any remaining unexpanded siblings of nodes along path
- Stop when node with goal state is expanded
- Include check that generated state has not already been explored along a path – cycle checking

Depth-First Search

```
def search(self):
    """Returns (next) path from the start node to a goal node.
    Returns None if no path exists.
    """
    while not self.empty_frontier():
        path = self.frontier.pop()
        self.num_expanded += 1
        if self.problem.is_goal(path.end()): # solution found
            self.solution = path # store solution
            return path
        else:
            neighs = self.problem.neighbors(path.end())
            for arc in reversed(neighs):
                self.add_to_frontier(Path(path,arc))
    # No more solutions
```

Depth-First Search




Depth-First Search — Analysis

- Storage: $O(bm)$ nodes (where m = maximum depth of search tree)
- Time: $O(b^m)$
- In cases where problem has many solutions, depth-first search may outperform breadth-first search because there is a good chance it will find a solution after exploring only a small part of the space
- However, depth-first search may get stuck following a deep or infinite path even when a solution exists at a relatively shallow level
- Therefore, depth-first search is not complete and not optimal
- Avoid depth-first search for problems with deep or infinite paths

Depth-Limited Search

- **Idea:** Impose bound on depth of a path
- In some problems you may know that a solution should be found within a certain cost (e.g. a certain number of moves) and therefore there is no need to search paths beyond this point for a solution
- Analysis
 - ▶ Complete but not optimal (may not find shortest solution)
 - ▶ However, if the depth limit chosen is too small a solution may not be found and depth-limited search is incomplete in this case
 - ▶ Time and space complexity similar to depth-first search (but relative to depth limit rather than maximum depth)

Iterative Deepening Search



Iterative Deepening Search

- It can be very difficult to decide upon a depth limit for search
- The maximum path cost between any two nodes is known as the **diameter** of the state space
- This would be a good candidate for a depth limit but it may be difficult to determine in advance
- **Idea:** Try all possible depth limits in turn
- Combines benefits of depth-first and breadth-first search

Iterative Deepening Search — Analysis

- Optimal; Complete; Space — $O(bd)$
- Some states are expanded multiple times: Isn't this wasteful?
 - ▶ Number of expansions to depth $d = 1 + b + b^2 + b^3 + \dots + b^d$
 - ▶ Therefore, for iterative deepening, total expansions = $(d+1)1 + (d)b + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + b^d$
 - ▶ The higher the branching factor, the lower the overhead (even for $b = 2$, search takes about twice as long)
 - ▶ Hence time complexity still $O(b^d)$
- Can double depth limit at each iteration – overhead $O(d \log d)$
- In general, iterative deepening is the preferred search strategy for a large search space where depth of solution is not known


Bidirectional Search

- **Idea:** Search forward from initial state and backward from goal state at the same time until the two meet
- To search backwards we need to generate predecessors of states (this is not always possible or easy)
- If operators reversible, successor sets and predecessor sets are the same
- If there are many goal states, maybe multi-state search would work (but not in chess)
- Need to check whether a node occurs in both searches – can be inefficient
- Which is the best search strategy for each half?

Bidirectional Search — Analysis

- If solution exists at depth d then bidirectional search requires time $O(2b^{\frac{d}{2}}) = O(b^{\frac{d}{2}})$ (assuming constant time checking of intersection)
- To check for intersection must have all states from one of the searches in memory, therefore space complexity is $O(b^{\frac{d}{2}})$

Bidirectional Search



Summary — Blind Search

Criterion	Breadth First	Uniform Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional
Time	b^d	b^d	b^m	b^l	b^d	$b^{\frac{d}{2}}$
Space	b^d	b^d	bm	bl	bd	$b^{\frac{d}{2}}$
Optimal	Yes	Yes	No	No	Yes	Yes
Complete	Yes	Yes	No	Yes, if $l \geq d$	Yes	Yes

b — branching factor

d — depth of shallowest solution

m — maximum depth of tree

l — depth limit